

Predicting Mispriced Polymarket Contracts Using Machine Learning

Logan Morof
CMSE 492 Final Project

December 7, 2025

Abstract

Prediction markets are supposed to aggregate information into prices that reflect the probability of real-world events, but in practice those prices wobble because liquidity is uneven, news is absorbed at different speeds, and traders behave imperfectly. In this project I build a supervised machine learning pipeline to flag Polymarket snapshots that later resolve far from their last traded price. I start with a small set of resolved markets to prove out feature engineering, then expand the dataset as the pipeline stabilizes. The task is to detect a mispricing label defined as a 15-cent gap between the final resolution price and the last observed price. Because only 3.7% of snapshots are mispriced, evaluation focuses on precision, recall, F1, PR-AUC, calibration, and threshold behavior rather than accuracy. A tuned XGBoost model achieves the best balance (ROC-AUC around 0.94, F1 around 0.59 at a 0.50 threshold). SHAP interpretation shows that volatility, order-flow pressure, and time to resolution are key drivers, and a simple backtest illustrates how different probability thresholds trade off coverage versus confidence. The results show that machine learning can surface short-lived inefficiencies, while also highlighting the practical hurdles of deploying such signals in live markets.

1 Background and Motivation

Polymarket lets traders buy YES/NO contracts on political, economic, and cultural events. In an efficient world, the price of a YES contract would track the probability of a “Yes” outcome. Real markets drift away from that ideal. Liquidity is lumpy and shallow in many contracts, news flows unevenly to different participants, and attention waxes and wanes. During those gaps, prices can be wrong for hours or days. Human intuition alone struggles to spot rare, transient mispricings consistently, especially under heavy class imbalance; a probabilistic model can impose a consistent scoring rule and surface weak signals that are easy to overlook. I initially treated this as an outcome prediction problem—guessing which side would win—but that framing felt unsatisfying for trading: many outcomes are obvious, imbalance is extreme, and a small edge is hard to monetize. After early tests, I pivoted to detecting mispricings near resolution: snapshots where the last traded price disagrees by at least fifteen cents with the eventual settlement. That reframing keeps the focus on short-lived, economically meaningful inefficiencies.

2 Machine Learning Task and Objective

The task is binary classification on market snapshots. A snapshot is labeled mispriced if the final resolution price diverges from the last observed price by at least fifteen cents:

$$y_{\text{misprice}} = \begin{cases} 1 & \text{if } |\text{final_price} - \text{last_price}| \geq 0.15, \\ 0 & \text{otherwise.} \end{cases}$$

Only 3.7% of snapshots satisfy this criterion. The goal is to produce calibrated probabilities that separate rare mispricings from the dominant normal cases. Because accuracy rewards trivial “always no” predictions, I rely on precision, recall, F1, ROC-AUC, PR-AUC, and calibration, plus an explicit look at threshold choices.

3 Data Collection and Description

All data come from Polymarket’s public Gamma, CLOB price history, and data APIs. I began with a small set of resolved markets to prototype feature engineering and model training, then expanded the universe once the pipeline was stable, pulling more price histories and trades. Each of the 7,558 snapshots captures a market at a specific time offset before resolution (12h, 6h, 3h, or 1h), combining price, order-flow, and timing context. In total there are 34 numeric features spanning rolling averages, moving volatility, price ranges, deviations from trend, buy/sell volumes, net order flow, trade counts, and timing variables such as age and time to resolution. After cleaning, there are no missing values (see `missingness_summary.csv`). The class balance is extremely skewed—279 positives and 7,279 negatives—shown in the updated class balance plots.

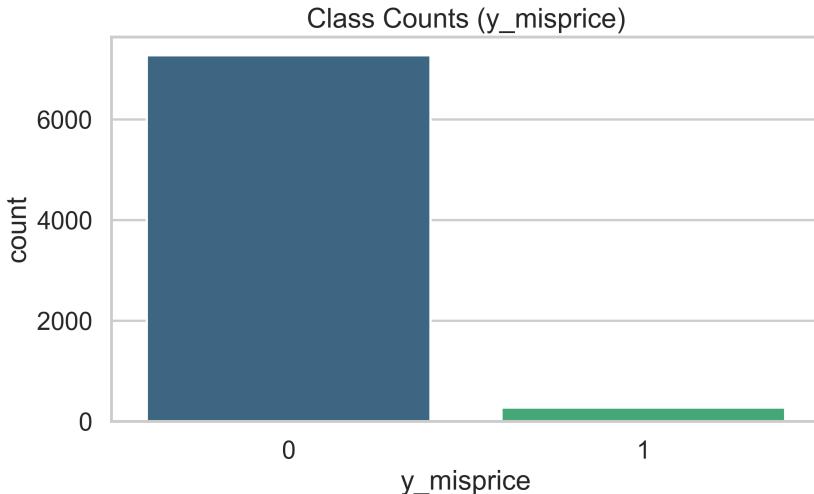


Figure 1: Class distribution of mispriced versus non-mispriced snapshots; only about 3.7% are labeled mispriced.

Univariate distributions and bivariate plots (see `univariate_summary_stats.csv` and the figures in `figures/`) show how key features behave overall and by label. Scatter plots of last versus final price and of time-to-resolution versus the absolute price gap highlight that mispriced points cluster where gaps are larger. A correlation heatmap (`correlation_heatmap.png`, `correlation_matrix.csv`) summarizes relationships among numeric features.

4 Preprocessing

I recompute y_{misprice} from the final and last prices for reproducibility. To avoid leakage, I drop `final_price` and any derived price gap used only for plotting. Features remain in natural units; tree-based models do not need scaling. The split is an 80/20 stratified train/test because no chronological split is available; class imbalance is handled with class weights and threshold tuning rather than oversampling. Although some descriptive plots use the full dataset for context, hyperparameter tuning and model selection are performed exclusively on the training split.

5 Models and Training Methodology

I move from simple to expressive models. Logistic regression provides a linear, class-weighted baseline. Random forest adds non-linear interactions via bagged trees. XGBoost pushes further with gradient boosting and tuned hyperparameters to capture subtle patterns in noisy, imbalanced data. All models minimize binary cross-entropy. Logistic regression needs minimal tuning; the forest and XGBoost use randomized search over depth, learning rate, subsampling, and regularization. Learning curves show the linear model plateauing quickly, the forest gaining modestly with more data, and XGBoost continuing to improve while keeping the gap between train and validation F1 reasonable.

Table 1: Model characteristics.

Model	Key hyperparameters	Loss	Regularization
Logistic Regression	C (class-weighted), lbfgs	BCE	ℓ_2
Random Forest	500 trees, depth=None, class_weight=balanced	Gini	depth/leaf controls
XGBoost (tuned)	600 trees, depth=6, lr=0.05, subsample=0.7, scale_pos_weight	BCE	ℓ_1, ℓ_2 , subsampling

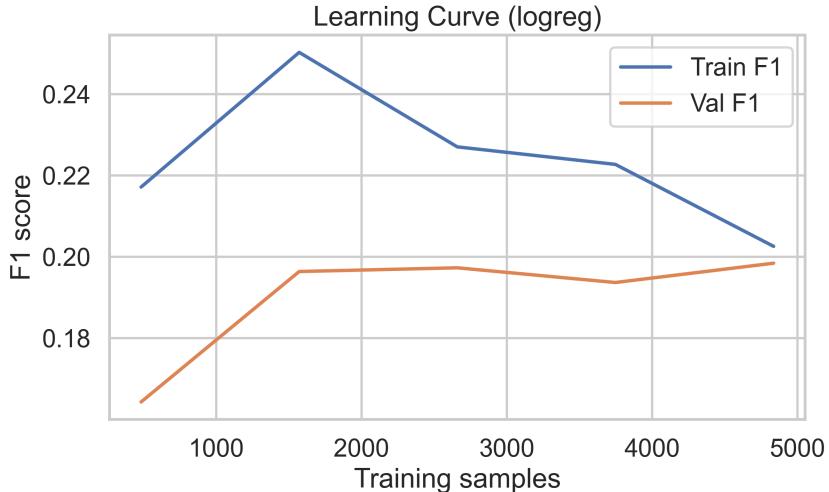


Figure 2: Learning curve for logistic regression. Validation F1 flattens quickly, indicating limited capacity for nonlinear structure.

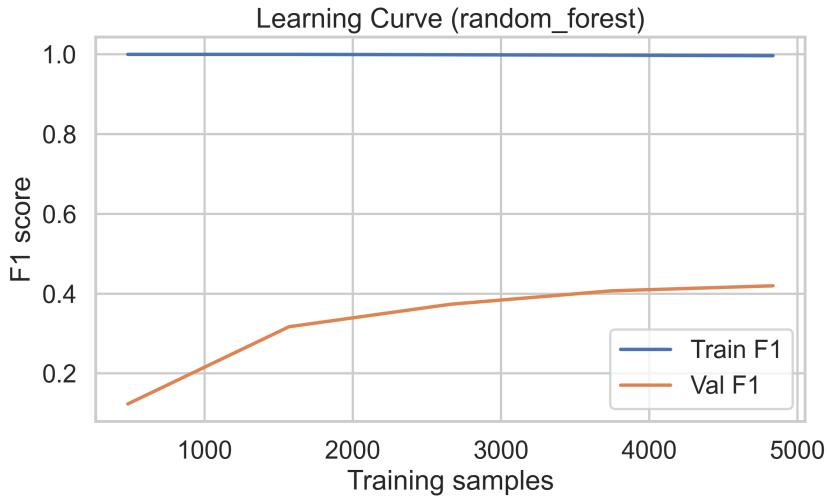


Figure 3: Learning curve for random forest. Nonlinear interactions help, but gains taper as data increase.

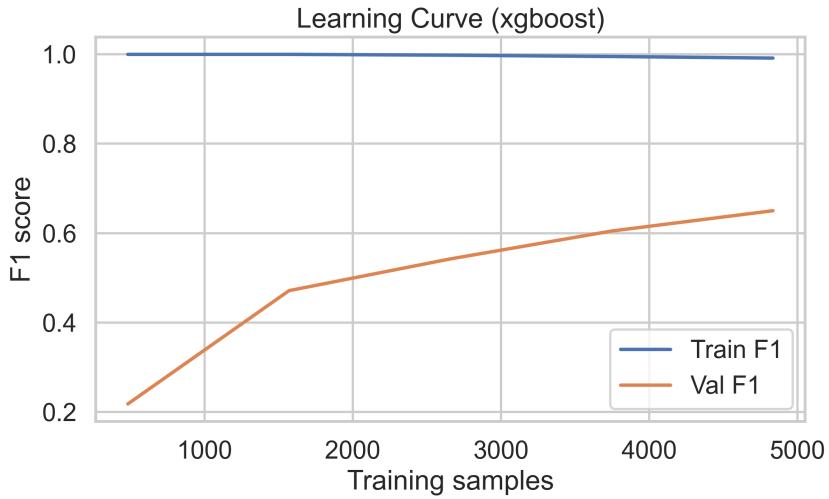


Figure 4: Learning curve for tuned XGBoost. Training and validation F1 rise together, suggesting better generalization without severe overfitting.

Training and inference remain fast. Logistic regression is effectively instantaneous; the forest and XGBoost train in about a second and predict in milliseconds per thousand samples.

Table 2: Training and inference time (`model_runtime_comparison.csv`).

Model	Train time (s)	Infer ms / 1k samples	Notes
Logistic Regression	0.042	0.0016	baseline-ish
Random Forest	1.223	0.0531	baseline-ish
XGBoost (tuned)	1.148	0.0089	optimized hyperparams

6 Metrics

With 3.7% positives, accuracy misleads. Precision answers how many flagged cases are real; recall measures how many real mispricings are caught; F1 balances the two. ROC-AUC summarizes ranking quality; PR-AUC focuses on rare positives; calibration checks probability reliability. Held-out metrics at the default 0.50 threshold are:

Table 3: Held-out metrics (`model_metrics_comparison.csv`).

Model	Accuracy	Precision	Recall	F1	ROC-AUC	PR-AUC
Logistic Regression	0.7586	0.1089	0.7679	0.1907	0.8383	0.1832
Random Forest	0.9656	0.5667	0.3036	0.3953	0.9433	0.5064
XGBoost (tuned)	0.9689	0.5763	0.6071	0.5913	0.9436	0.5742

7 Results and Model Comparison

The ROC curve for tuned XGBoost stays well above the diagonal, and the PR curve stays above the baseline implied by 3.7% positives, indicating good ranking of rare positives.

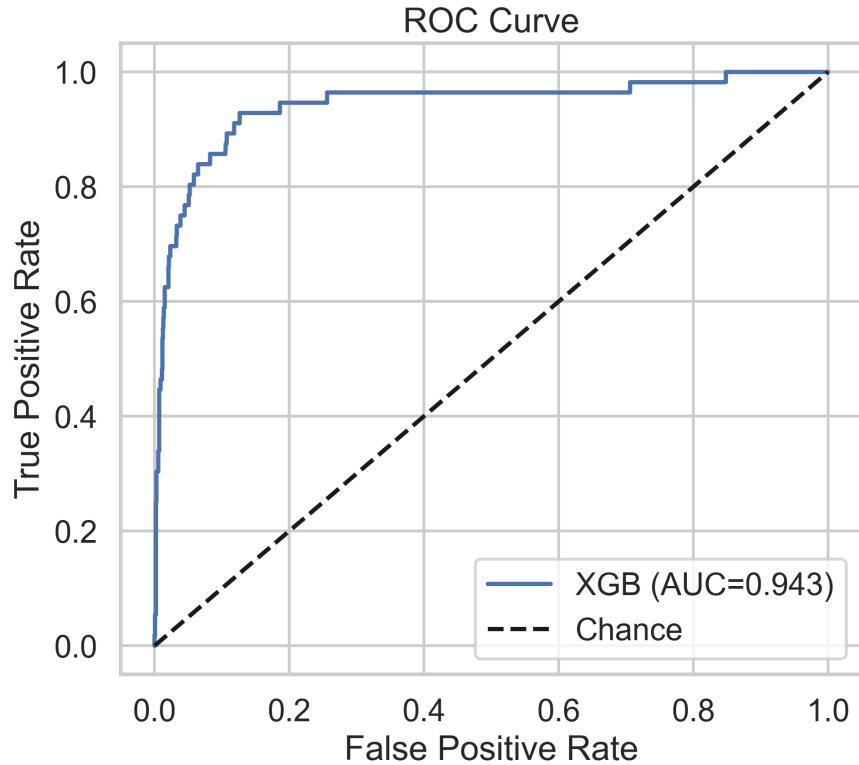


Figure 5: ROC curve for tuned XGBoost; ROC-AUC is about 0.94, meaning true mispricings generally score higher than normal cases.

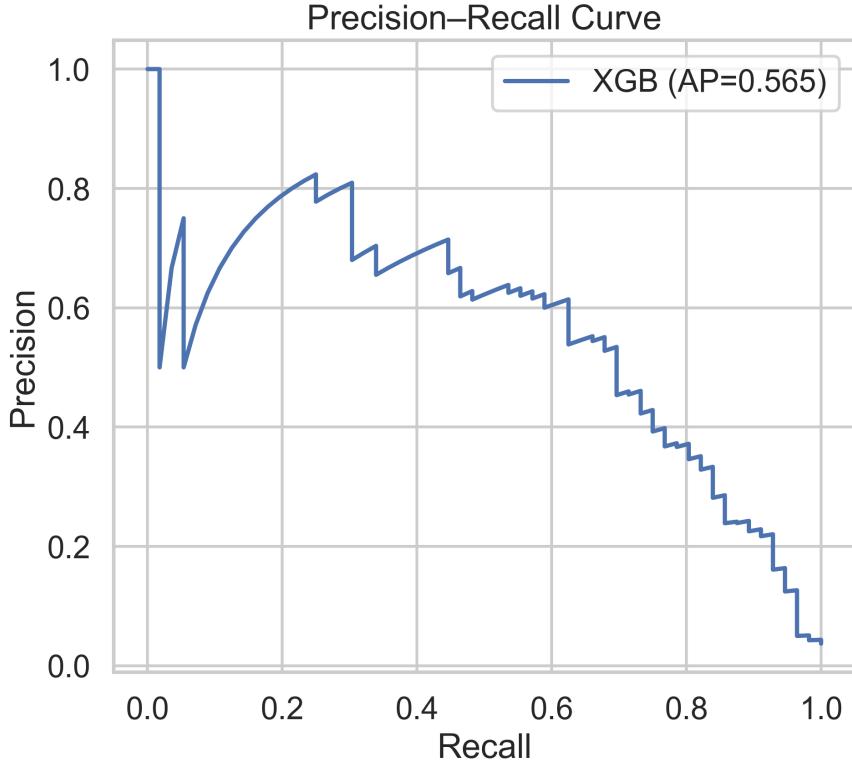


Figure 6: Precision–Recall curve; PR-AUC around 0.57 shows better-than-chance precision even as recall rises under heavy imbalance.

Threshold sweeps show how moving the cutoff trades recall for precision. Around 0.50, F1 peaks; around 0.70, precision improves at the cost of recall.

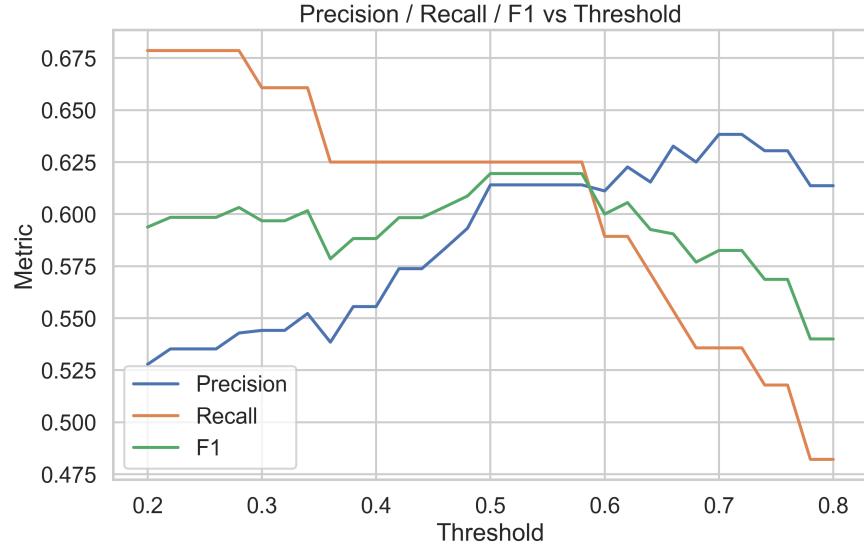


Figure 7: Threshold sweep. Lower thresholds catch more mispricings with more false alarms; higher thresholds are more selective and precise.

Confusion matrices make the trade-offs concrete. At $t = 0.50$, the model finds 35 true mispricings, misses 21, and produces 22 false alarms out of 1,456 normal cases ($TN=1434$, $FP=22$, $FN=21$, $TP=35$). At $t = 0.70$, it finds 30, misses 26, and false alarms drop to 17 ($TN=1439$, $FP=17$, $FN=26$, $TP=30$).

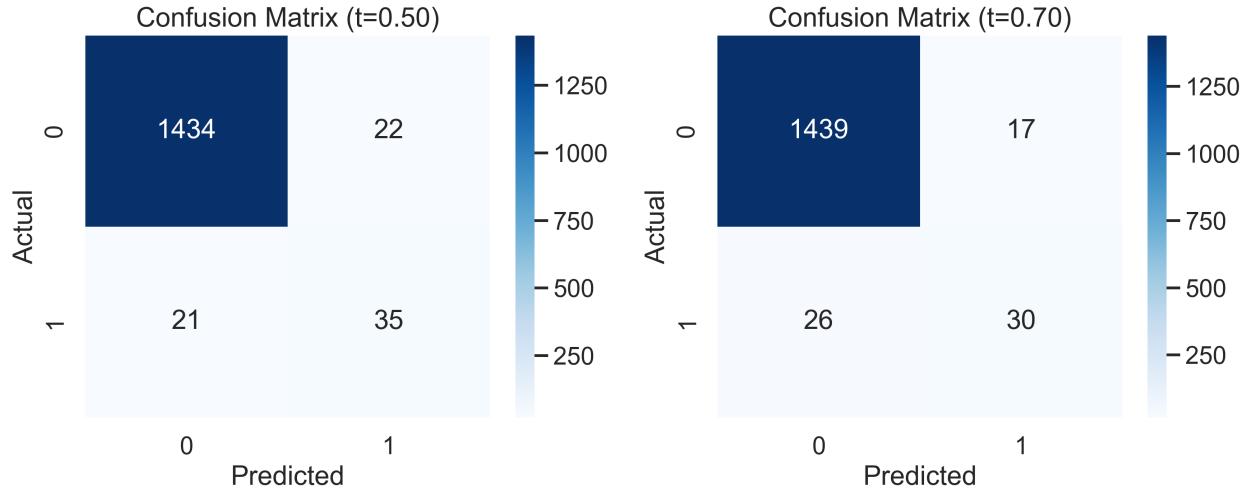


Figure 8: Confusion matrices for $t = 0.50$ (left) and $t = 0.70$ (right). Higher thresholds reduce false positives but miss more true mispricings.

Calibration curves show slight underconfidence in mid-range probabilities: a predicted 40% corresponds to an empirical frequency somewhat higher. That conservatism is acceptable for ranking, but a post-hoc calibration step could help if probabilities drive execution.

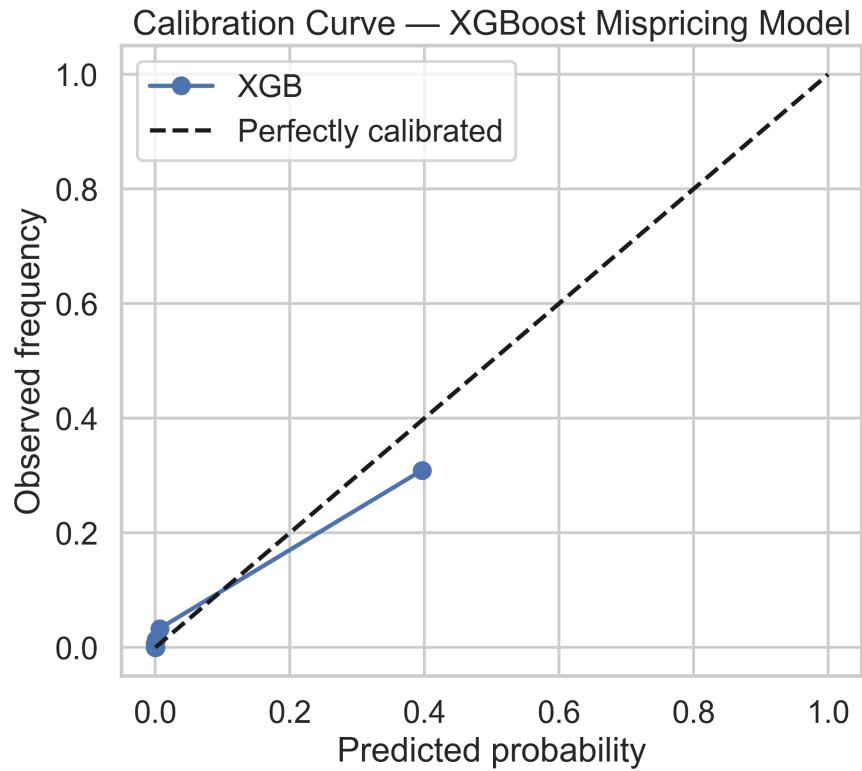


Figure 9: Calibration curve. Mid-range probabilities are conservative; a calibration step could refine them for trading.

A cost-sensitive view echoes the threshold sweep: if missing a mispricing is worse than a false alarm, lower thresholds are favored; if false alarms are costly, higher thresholds are better.

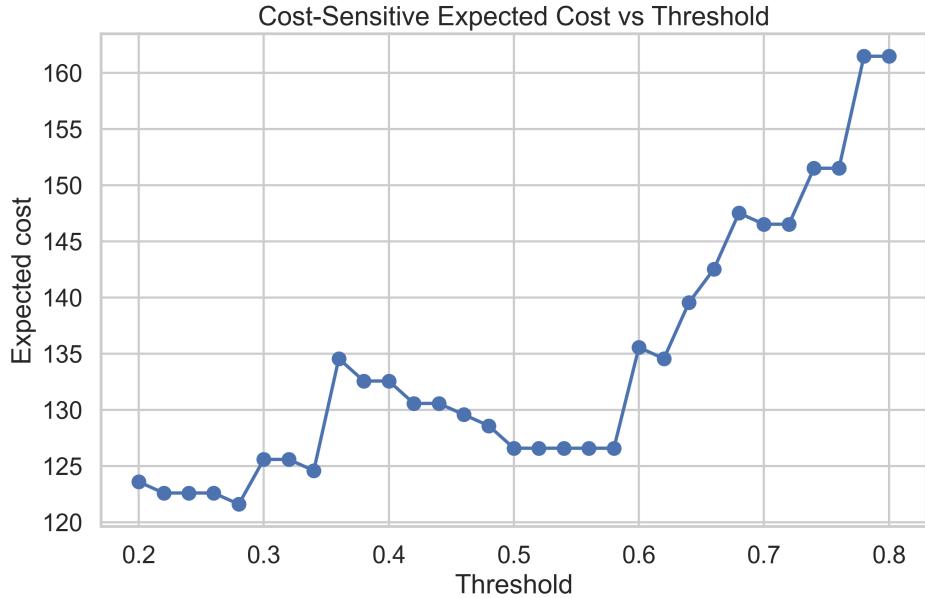


Figure 10: Cost-sensitive expected cost versus threshold. The preferred threshold shifts with the relative cost of misses versus false alarms.

8 Model Interpretation

Feature importances from XGBoost show that short-term volatility, deviations from recent averages, asymmetric order flow, and proximity to resolution dominate the signal. That matches intuition: stressed markets near the finish line are prone to temporary mispricing.

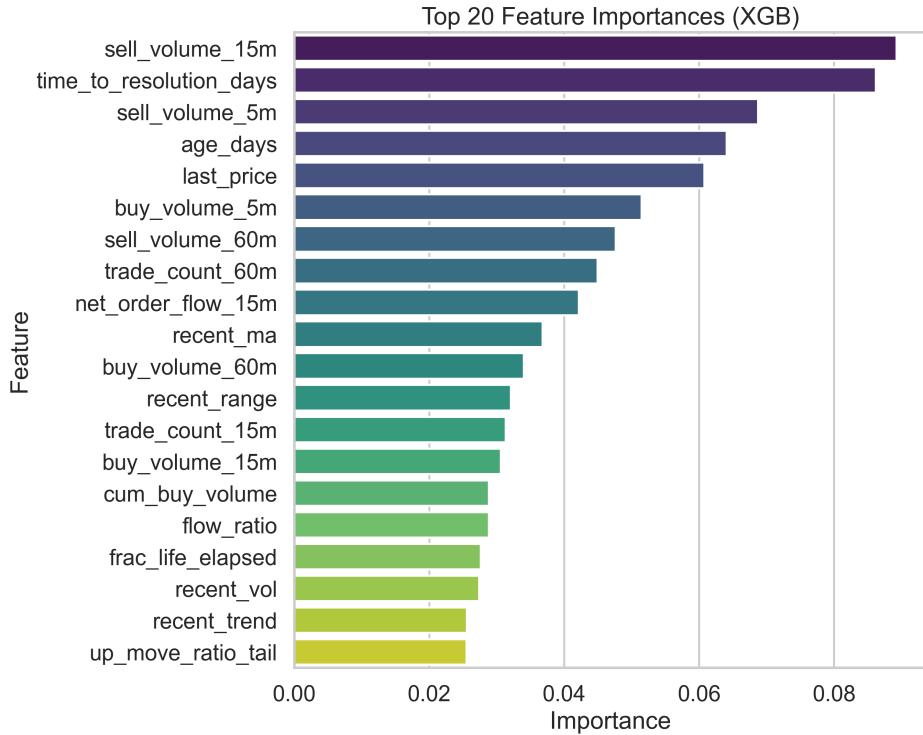


Figure 11: Top XGBoost feature importances. Volatility, order-flow pressure, and time to resolution are key drivers.

SHAP values explain predictions at the instance level. The summary plot shows which features push probabilities up or down across the dataset, while dependence plots reveal interactions. Heavy recent trading and short time to resolution raise mispricing risk; calm, liquid markets tilt toward fair pricing.

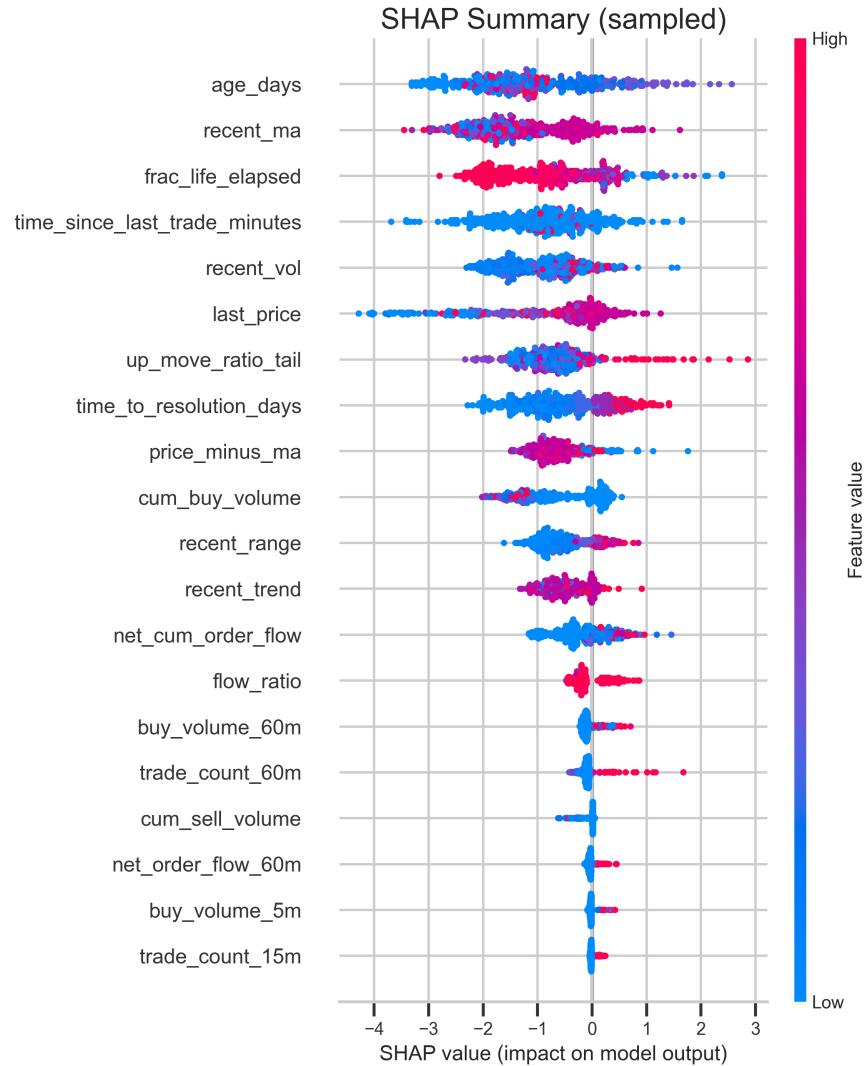


Figure 12: SHAP summary. Wide spreads indicate strong influence; volatility and order-flow asymmetry tend to push predictions toward mispriced.

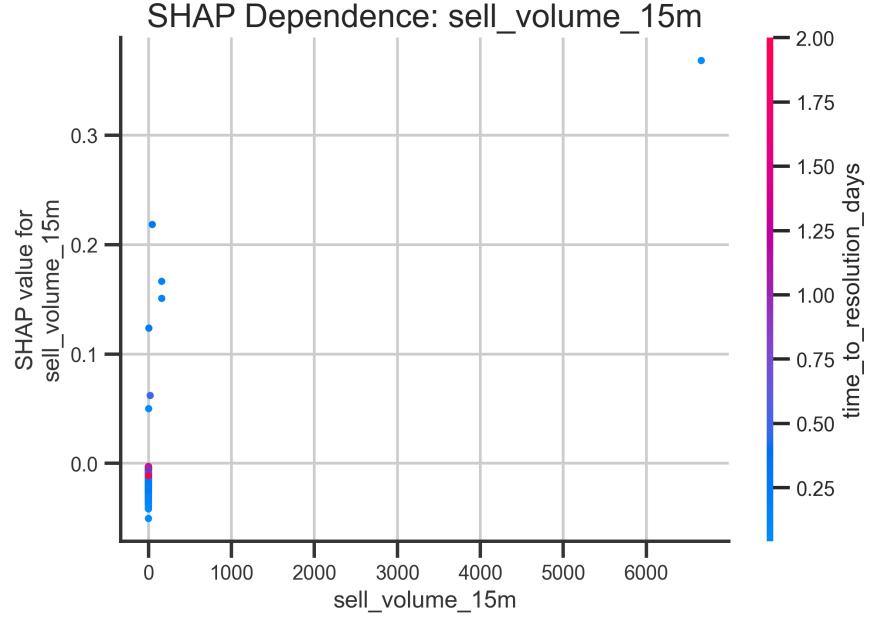


Figure 13: SHAP dependence for fifteen-minute sell volume. Sharp bursts of selling, especially near resolution, increase mispricing risk.

9 Backtesting Analysis

To connect scores to trading implications, I ran a simple optimistic backtest on the held-out set. Per-trade PnL is the absolute gap between final and last price; fees, slippage, and order book depth are ignored. Three strategies are compared: an unattainable oracle that trades only on true mispricings, a balanced 0.50 threshold strategy, and a more selective 0.70 threshold strategy. The oracle earns the most. The 0.50 strategy executes more often with moderate gains; the 0.70 strategy trades less but with higher confidence and steadier outcomes.

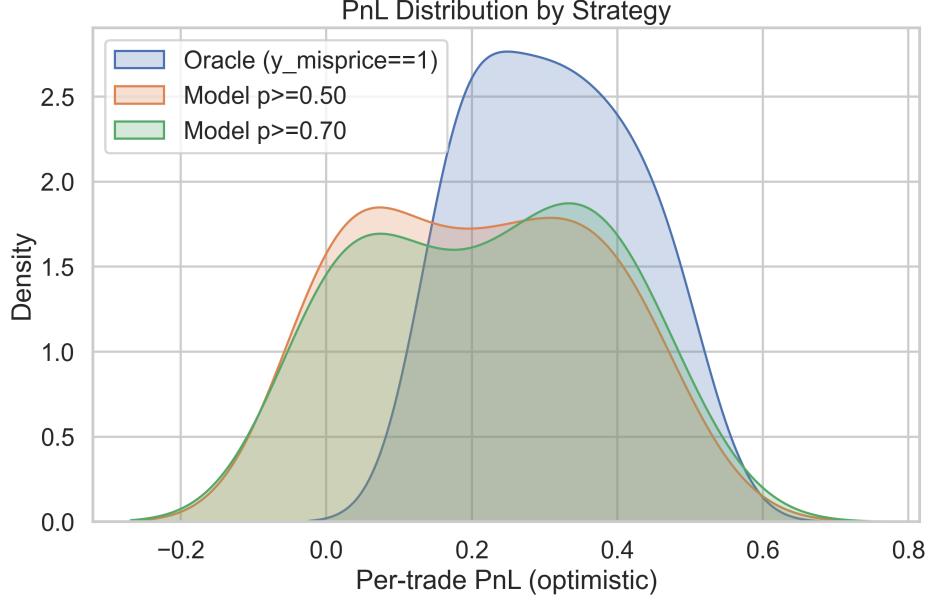


Figure 14: Per-trade PnL distributions. The 0.70 strategy trades less often but with a tighter, more favorable distribution; the 0.50 strategy trades more frequently with a broader spread.

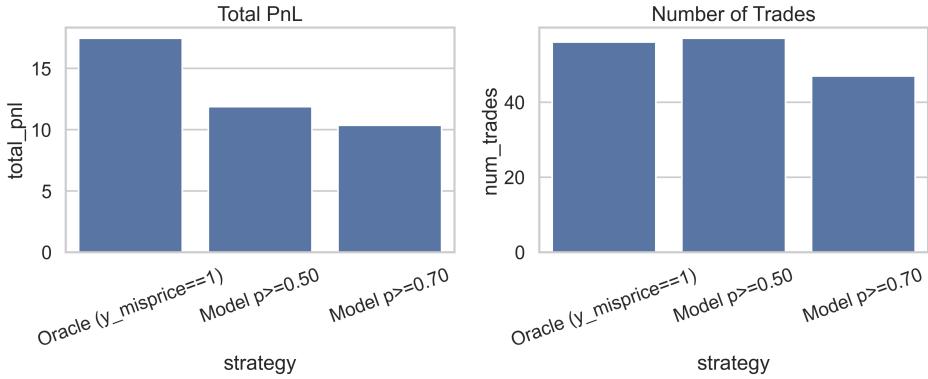


Figure 15: Total PnL and trade counts. The conservative threshold sacrifices volume for precision; the balanced threshold captures more opportunities with more noise.

10 Conclusion

Reframing from outcome prediction to mispricing detection produced a model that surfaces short-lived inefficiencies in Polymarket. A tuned XGBoost classifier delivers the best balance of precision and recall under heavy imbalance, with ROC-AUC around 0.94 and F1 around 0.59 at a 0.50 threshold. Interpretation links the signal to intuitive dynamics: volatility spikes, asymmetric order flow, and looming resolution all coincide with higher mispricing risk. Threshold tuning and a simple backtest show how to trade off coverage against confidence.

Limitations remain. The backtest is optimistic and omits execution costs, slippage, and depth; the split is stratified rather than chronological, so live deployment would need walk-forward validation; and market regimes change, so models require monitoring and retraining. Future extensions

include rolling validation, richer on-chain or liquidity features, and a live paper-trading loop to observe how calibrated probabilities and thresholds behave in real time.

References

- [1] Polymarket api reference. <https://data-api.polymarket.com>. Accessed 2025-01-10.
- [2] Polymarket documentation. <https://docs.polymarket.com/>. Accessed 2025-01-10.
- [3] Tianqi Chen and Carlos Guestrin. Xgboost documentation. <https://xgboost.readthedocs.io/>, 2016.
- [4] Pedregosa et al. Scikit-learn machine learning in python. <https://scikit-learn.org/stable/>, 2011.
- [5] J. D. Hunter and Matplotlib Developers. Matplotlib: Visualization with python. <https://matplotlib.org>, 2024.
- [6] Scott M. Lundberg and Su-In Lee. Shap: Shapley additive explanations. <https://shap.readthedocs.io/>, 2017.
- [7] OpenAI. Chatgpt (gpt-5.1 model). <https://chat.openai.com>, 2025. Used extensively for drafting, debugging, analysis, and project collaboration.
- [8] OpenAI. Github copilot / openai codex. <https://github.com/features/copilot>, 2025. Used for code generation and notebook construction assistance.
- [9] The pandas development team. pandas: Python data analysis library. <https://pandas.pydata.org/>, 2024.
- [10] Michael Waskom. Seaborn statistical visualization. <https://seaborn.pydata.org>, 2024.